

DTU



Perception for Autonomous Systems 31392:

Visual SLAM

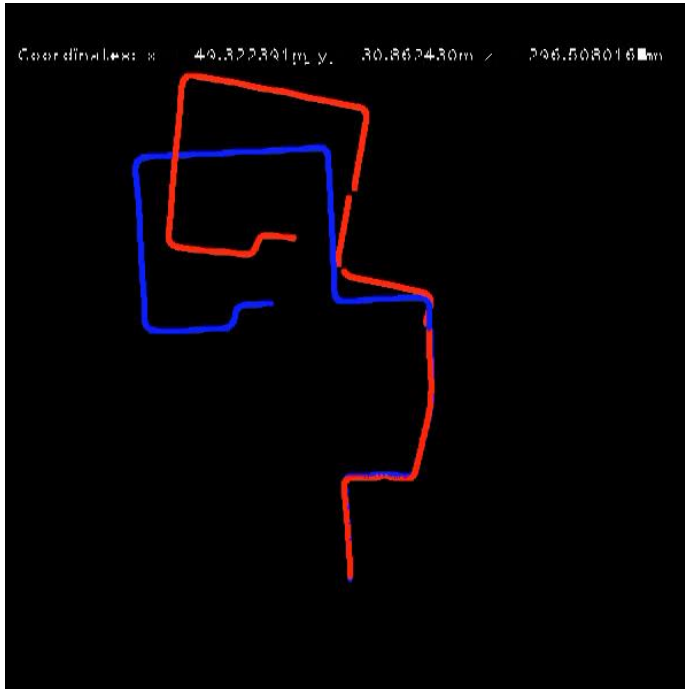
Simultaneous Localization & Mapping

Lecturer: Evangelos Boukas—PhD

- Sum up Localization from last time
- Some terminology
- Pose-Landmark Graph Slam
- Example of Linear 1D SLAM
- Non-Linear Optimization approaches
- Bundle Adjustment
- Visual Slam System architecture
- ORBSLAM

Visual Odometry is great


- Lets Sum Up What we did last time
- We performed 3D-to-2D



Visual Odometry is great

- Lets Sum Up What we did last time
- We performed 3D-to-2D

Coordinates: x: 53.169407m, y: 1.180850m, z: 0.37474980m






Visual Odometry is great

- Anything more we mentioned?
 - Windowed Bundle Adjustment (BA)

Windowed Bundle Adjustment (BA)

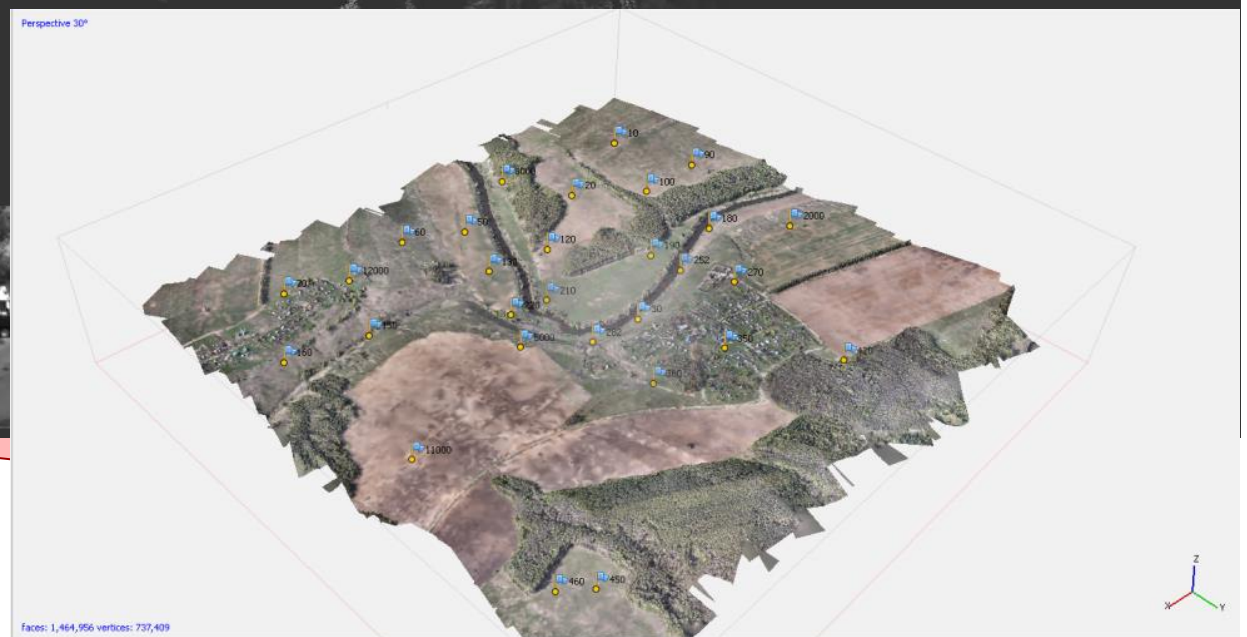



- Similar to pose-optimization but it also optimizes 3D points ^{m}
- In order to not get stuck in local minima, the initialization should be close the minimum
- Levenberg-Marquadt can be used

Formal definitions

- Visual Odometry
- Structure from Motion (SfM)
- Bundle Adjustment
- Visual SLAM

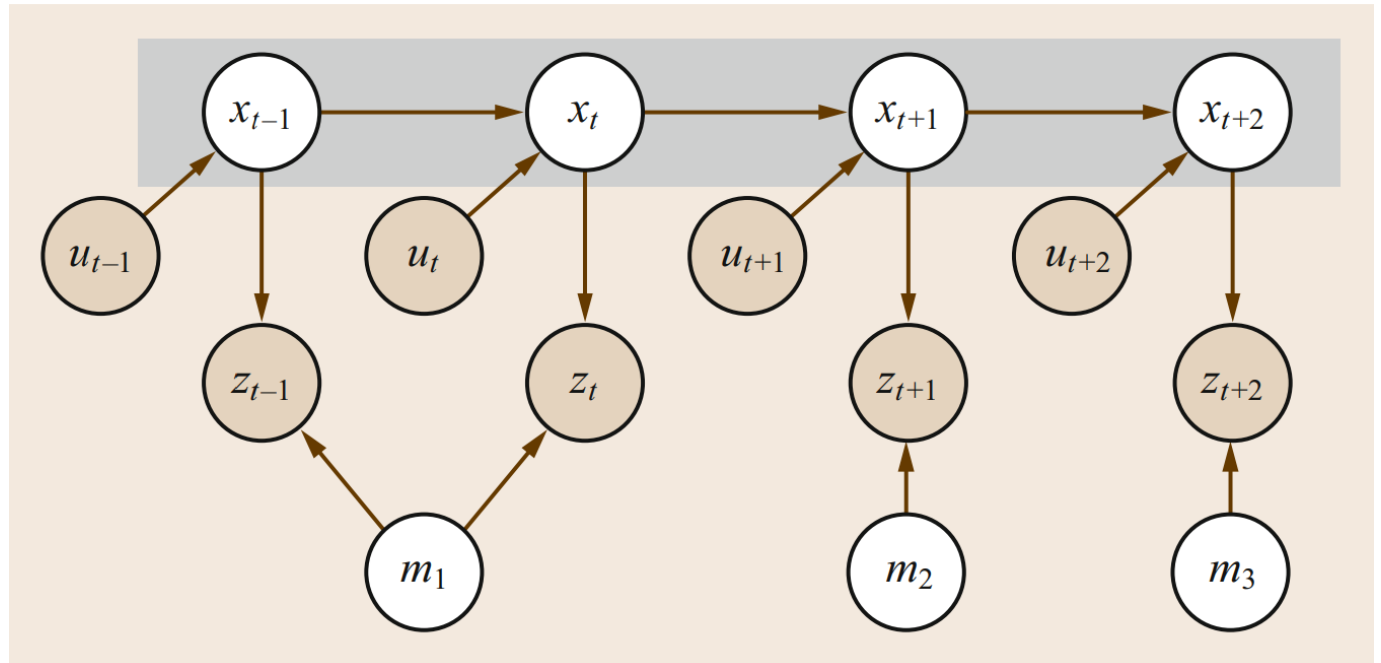
KITTI 00 (Full SLAM)



- Sum up Localization from last time
- Some terminology
- Pose-Landmark Graph Slam
- **Example of Linear 1D SLAM**
- Non-Linear Optimization approaches
- Bundle Adjustment
- Visual Slam System architecture
- ORBSLAM

Pose-Landmark Graph-Slam

- SLAM problem depicted as Bayes network graph
- At each **location** \mathbf{x}_t
- Observes a nearby feature in the **map** $\mathbf{m} = \{\mathbf{m}_1; \mathbf{m}_2; \mathbf{m}_3\}$
- Movement \mathbf{u}_t
- An arrow defines causal relationship





Graph-Based SLAM

Definition

- Use a graph to represent the problem
- Nodes represent:
 - poses or
 - locations
- Edges Represent:
 - Landmark observations
 - Odometry Measurements
- The minimization optimizes the landmark locations and robot poses

Graph-Based SLAM: Build the graph and find a node configuration that minimize the error introduced by the constraints




Graph-Based SLAM (intuition of optimization)

- Observing previously seen areas generates constraints between non-successive poses
- Treat constraints (motion and measurement) as “soft” elastic springs
- Want to minimize the total energy in the springs

We can define the error as follows

- Expected observation (2D sensor)

– With the error: $e_{ij}(\mathbf{x}_i, \mathbf{x}_j) = \hat{\mathbf{z}} - \mathbf{I}$



1D Linear SLAM

- In the linear case we can solve as follows:

- First construct all constrains

- Absolute Constrains:

$$X(0) = Q - \text{starting position}$$

- Movement Constrains:

$$X(t) = X(t-1) + Dx(t)$$

- Measurement constrains:

$$L(k) = X(t) + N$$

- Then, solve linear equations



X0



X1



X2



X3



L0

1D Linear SLAM – case 1

- Case 1 - Exact solution exists:


$$\begin{bmatrix} & & \\ & & \\ & & \end{bmatrix} \cdot \begin{bmatrix} \\ \\ \end{bmatrix} = \begin{bmatrix} \\ \\ \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & -1 & 1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \end{bmatrix}$$

$$A * X = B \quad X = A^{-1} * B$$


$$A^{-1} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

$$x = \begin{bmatrix} -3 & 2 & 5 \end{bmatrix}$$



1D Linear SLAM – case 2

- Case 2 – Overdefined problem:
 - X_0 sees L_0 at distance 10
 - X_1 sees L_0 at distance 5
 - X_2 sees L_0 at distance 2



$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ 0 & -1 & 1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$

$$A \cdot X = B \quad X = A^{-1} \cdot B \quad A^{-1} = [?]$$

$$X = (A^T \cdot A)^{-1} \cdot A^T \cdot B$$


$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ 0 & -1 & 1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$

$$x = \begin{bmatrix} -3 & 2 & 5 & 7 \end{bmatrix}$$

We infer a consistent landmark position

1D Linear SLAM – case 3

- Case 3 – Inconsistent Measurements :
 - X_0 sees L_0 at distance 10
 - X_1 sees L_0 at distance 5
 - X_2 sees L_0 at distance 1 (**Wrong**)



$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 1 & 0 & 0 \\ 0 & -1 & 1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \cdot \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L_0 \end{bmatrix} = \begin{bmatrix} -3 \\ 5 \\ 3 \\ -10 \\ -5 \\ -1 \end{bmatrix}$$

$$X = (A^T * A)^{-1} * A^T * B$$

$$x = \begin{bmatrix} -3 & 2.125 & 5.5 & 6.875 \end{bmatrix}$$

We handled inaccurate measurements

1D Linear SLAM – case 4


Case 4 – Inconsistent Measurements with Confidence Matrix:

- Linear Least Squares allows us to include a weighting of each linear constraint.
- We can include weights in the computation
- We weight each constraint by a diagonal matrix where the weights are $1/\text{variance}$ for each constraint.
- Let's say X_2 **variance is 5**

$$X = (A^T * W * A)^{-1} * A^T * W * B$$

$$x = \begin{bmatrix} -3 & 2.18 & 5.71 & 6.82 \end{bmatrix}$$

Why did the estimation just become worse??



What about non-Linear Least Squares?

- Large number of geometric problems in computer vision are non-linear least-squares problems.

$$\mathbf{x} = \mathbf{h}(\theta)$$

where $\mathbf{h} : \mathbf{R}^n \rightarrow \mathbf{R}^m$.

- \mathbf{x} is the measurement vector, θ is the parameter vector.
- Write $\mathbf{f}(\theta) = \mathbf{h}(\theta) - \mathbf{x}$.
- We desire to minimize

$$\|\mathbf{f}(\theta)\|^2$$

over all choices of parameter θ .

Gauss Newton Solution

1. Start from an initial value θ_0 .
2. At step i assume a linear approximation for the function at θ_i

$$\mathbf{f}(\theta_i + \Delta) = \mathbf{f}(\theta_i) + \mathbf{f}_\theta \Delta \text{ where } \mathbf{f}_\theta = \partial \mathbf{f} / \partial \theta = \mathbf{J} .$$

3. Solve

$$\mathbf{f}(\theta_i + \Delta) = \mathbf{f}(\theta_i) + \mathbf{J} \Delta = 0$$

or

$$\mathbf{J} \Delta = -\mathbf{f}(\theta_i)$$


4. This is a linear least-squares problem (solve for Δ):

$$\mathbf{J}^\top \mathbf{J} \Delta = \mathbf{J}^\top \mathbf{f}(\theta_i)$$


5. Then set $\theta_{i+1} = \theta_i + \Delta$.

Gauss-Newton update equation

$$\mathbf{J}^\top \mathbf{J} \Delta = -\mathbf{J}^\top \mathbf{f}$$



1D Gauss-Newton (Newton)
iteration.



1D Gauss-Newton (Newton)
iteration (failure)


Gradient Descent

Search direction is the direction of fastest descent of the function g .

Gradient descent update equation

$$\lambda \Delta = -g_{\theta} = -\mathbf{J}^{\top} \mathbf{f}$$

Requires a 1D line search in λ to find the optimum direction.




Levenberg-Marquadt


- Mixture of Gauss-Newton and Gradient descent.
 - Acts like Gauss-Newton when close to the minimum (quadratic region)
 - Gradient descent when improvement is difficult.
 - Depends on a parameter λ which
 1. Controls the mixture of Gauss-Newton and Gradient Descent
 2. Controls the step-length.
-

What about non-Linear Least Squares?


- Lets See some examples 1:



Gauss-Newton




Gradient descent




Levenberg

What about non-Linear Least Squares?


- Lets See some examples 2:



Gauss-Newton




Gradient descent




Levenberg

What about non-Linear Least Squares?


- Lets See some examples 3:



Gauss-Newton



Gradient descent



Levenberg

- It is obvious that Levenberg Marquadt displays robustness

Bundle Adjustment

- Bundle Adjustment is the employment of nonlinear optimization in the problem of the minimization of the re-projection error, by finding the optimal Poses (extrinsics) of the cameras and the locations of the 3D points.


$$\arg \min_{\mathbf{w}, \boldsymbol{\theta}} \sum_{f=1}^F \sum_{n=1}^N ||\mathbf{x}_n^f - \pi(\mathbf{w}_n; \boldsymbol{\theta}^f)||_2^2$$

\mathbf{x} 2D projection \mathbf{w} 3D point

$\boldsymbol{\theta}$ extrinsics N no. of points

π projection function

F no. of frames



Bundle Adjustment – Linearization

$$\pi(\mathbf{w}_n + \Delta \mathbf{w}_n; \boldsymbol{\theta}_f \circ \Delta \boldsymbol{\theta}_f) \approx \pi(\mathbf{w}_n; \boldsymbol{\theta}_f) + \mathbf{J}_n^f \begin{bmatrix} \Delta \boldsymbol{\theta}_f \\ \Delta \mathbf{w}_n \end{bmatrix}$$



$$\arg \min_{\Delta \boldsymbol{\theta}, \Delta \mathbf{w}} \sum_{f=1}^F \sum_{n=1}^N \rho_n^f \left\| \mathbf{x}_n^f - \pi(\mathbf{w}_n; \boldsymbol{\theta}_f) - \mathbf{J}_n^f \begin{bmatrix} \Delta \boldsymbol{\theta}_f \\ \Delta \mathbf{w}_n \end{bmatrix} \right\|_2^2$$

\mathbf{x} 2D projection $\mathbf{w} \leftarrow$ 3D point

$\boldsymbol{\theta}$ extrinsics N no. of points

π projection function

F no. of frames $\rho \rightarrow$ visibility $\in [0, 1]$


Bundle Adjustment – Linearization

- The Linearization of the minimization happens by calculating the Jacobian of the projection matrix
- Assuming the following projection matrix:
$$\begin{bmatrix} wu \\ wv \\ w \end{bmatrix} = \begin{bmatrix} f_x & s_k & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} [R \quad T] \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$
- We first define the orientation as the rotation matrix associated with the axis angle w_x, w_y, w_z using the Rodrigues equation
- The Jacobian **FUNCTION** can be calculated as:

$$\mathbf{J} = \begin{bmatrix} \frac{\partial u}{\partial w_x} & \frac{\partial u}{\partial w_y} & \frac{\partial u}{\partial w_z} & \frac{\partial u}{\partial f} & \frac{\partial u}{\partial u_0} & \frac{\partial u}{\partial v_0} & \frac{\partial u}{\partial X} & \frac{\partial u}{\partial Y} & \frac{\partial u}{\partial Z} \\ \frac{\partial v}{\partial w_x} & \frac{\partial v}{\partial w_y} & \frac{\partial v}{\partial w_z} & \frac{\partial v}{\partial f} & \frac{\partial v}{\partial u_0} & \frac{\partial v}{\partial v_0} & \frac{\partial v}{\partial X} & \frac{\partial v}{\partial Y} & \frac{\partial v}{\partial Z} \end{bmatrix}$$

Bundle Adjustment – Comments

- Bundle adjustment (and graph optimization) is the backbone of all SLAM algorithms
 - Keep in mind that:
 - We need to provide the Jacobian of the projection
 - We usually provide a covariance matrix (See the linear case for uncertainty)
 - It is solved using Levenberg-Marquadt
 - There are a lot of computational issues which we can overcome by exploiting the sparsity of the function $AX=b$ (see least squares)
- Look at the following A and b Matrices
- This solution is called **Sparse Bundle Adjustment!**



Bundle Adjustment – Covariance


- The uncertainty of the camera pose C_k is a combination of the uncertainty at C_{k-1} (black-solid ellipse) and the uncertainty of the transformation T_k (gray dashed ellipse)

- $C_k = f(C_{k-1}, T_k)$

- The combined covariance Σ_k is


$$\begin{aligned}\Sigma_k &= J \begin{bmatrix} \Sigma_{k-1} & 0 \\ 0 & \Sigma_{k,k-1} \end{bmatrix} J^\top \\ &= J_{\vec{C}_{k-1}} \Sigma_{k-1} J_{\vec{C}_{k-1}}^\top + J_{\vec{T}_{k,k-1}} \Sigma_{k,k-1} J_{\vec{T}_{k,k-1}}^\top\end{aligned}$$

- The camera-pose uncertainty is always increasing when concatenating transformations. Thus, it is important to keep the uncertainties of the individual transformations small



Recent Visual Slam Solutions - Intro

- That was too much info, let's see now some recent solutions to the Slam Problem:
- Most recent visual SLAM methods are split in two parts:
- The **Frontend**: where the raw data are converted into pose graphs and Loop constraints and the **Backend** where, given a graph with constraints, the new pose of the robot is calculated as well as the surrounding map points.



Recent Visual Slam Solutions – Common Architecture

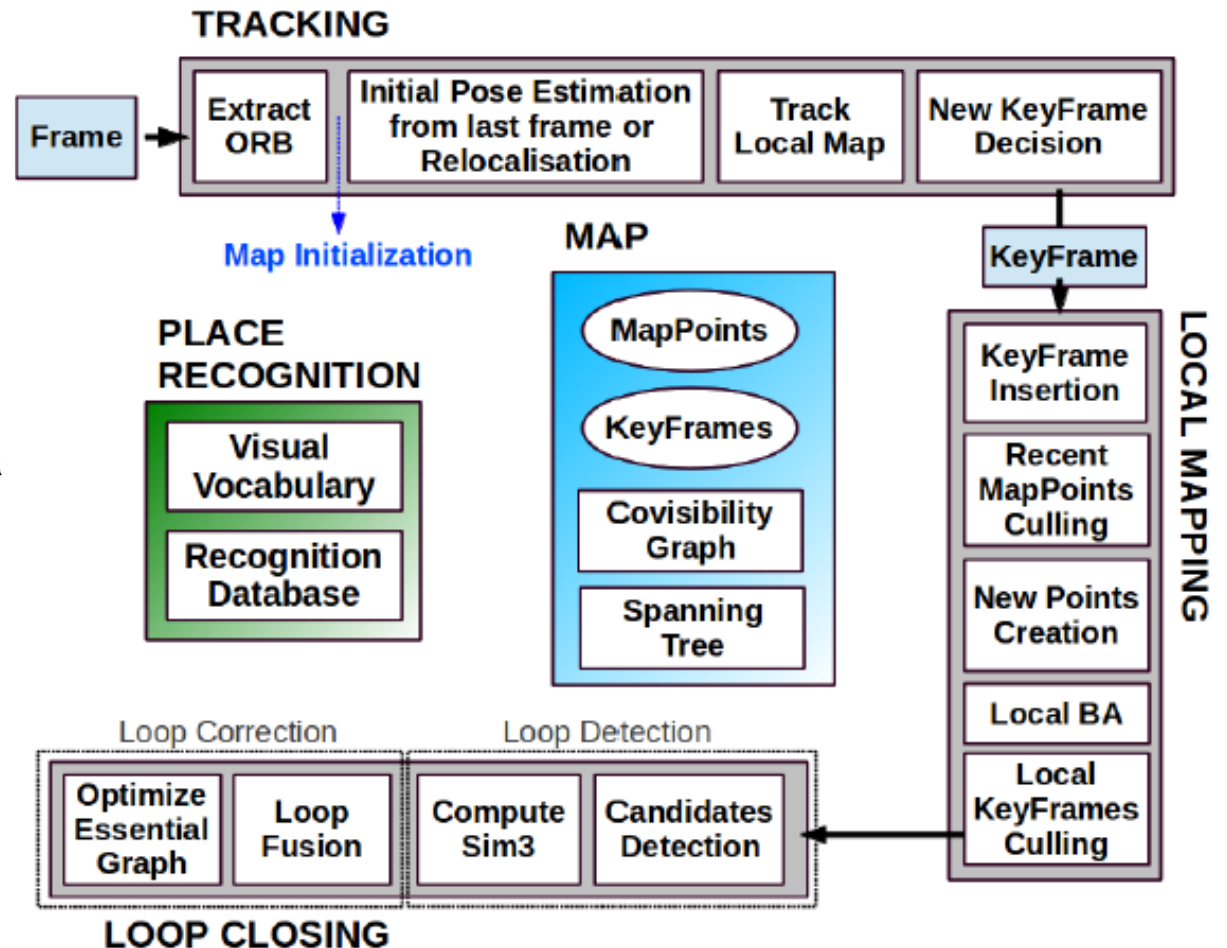
- Front End
 - Data Association
 - Frame to Frame
 - Multi-frame
 - Loop Closure Detection
 - Geometric Initialization
 - Pose Estimation
 - Landmark Triangulation
 - System Formation
 - Observation Matrix
 - Covariance Matrix
 - Graph Generation and Update
- Back End
 - Filter-Based State Estimation
 - Extended Kalman Filter
 - Particle Filters
 - Least squares optimization
 - Bundle Adjustment
 - Graph Optimization
 - Key Frame

Recent Visual Slam Solutions – recent advances

Towards Realtime operation?:

- The computational cost of bundle adjustment has lead to the idea of **keyframing**:
i.e.: identifying and describing some of the frames to be used for graph optimization.
- Bags of words for robust loop closure.
 - What can you tell me about that?
- Co-visibility Graph

- The ORBSLAM algorithm is one of the most well performing opensource implementations of visual slam.
- Three parallel threads:
 - tracking,
 - localizing the camera with every frame and deciding when to insert a new keyframe
 - local mapping
 - Processes new keyframes and performs local BA to achieve an optimal
- reconstruction in the surroundings of the camera
 - loop closing
 - The loop closing searches for loops with every new keyframe
 - Essential Graph



ORB SLAM on Kitty

ORB-SLAM

Raúl Mur-Artal, J. M. M. Montiel and Juan D. Tardós

{raulmur, josemari, tardos} @unizar.es



Instituto Universitario de Investigación
en Ingeniería de Aragón
Universidad Zaragoza



Universidad
Zaragoza

Sum up

- Sum up Localization from last time
- Some terminology
- Pose-Landmark Graph Slam
- Example of Linear 1D SLAM
- Non-Linear Optimization approaches
- Bundle Adjustment
- Visual Slam System architecture
- ORBSLAM

Perception for Autonomous Systems 31392:

Visual SLAM

Simultaneous Localization & Mapping

Lecturer: Evangelos Boukas—PhD